

EDA and Data Prediction on the Telcom Churn Dataset

"Churning" refers to the number of Customers or Employees that leave a Company in a given time period. Changes in a business's churn rate can provide valuable insight into an organization.

Importing librariers required for Data Analysis and Predictions

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: df=pd.read_csv("telecom_churn.csv")
```

```
In [3]: df.head()
```

Out[3]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DevicePro
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	

5 rows × 21 columns



```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   customerID            7043 non-null   object
 1   gender                7043 non-null   object
 2   SeniorCitizen         7043 non-null   int64
 3   Partner               7043 non-null   object
 4   Dependents            7043 non-null   object
 5   tenure               7043 non-null   int64
 6   PhoneService          7043 non-null   object
 7   MultipleLines         7043 non-null   object
 8   InternetService       7043 non-null   object
 9   OnlineSecurity        7043 non-null   object
10  OnlineBackup          7043 non-null   object
11  DeviceProtection      7043 non-null   object
12  TechSupport           7043 non-null   object
13  StreamingTV           7043 non-null   object
14  StreamingMovies       7043 non-null   object
15  Contract              7043 non-null   object
16  PaperlessBilling      7043 non-null   object
17  PaymentMethod         7043 non-null   object
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   object
20  Churn                 7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

Handling Null values and then replacing them with Central Tendencies

```
In [5]: df[df["TotalCharges"]==" "]
#TotalCharges column has 11 rows without any value
```

Out[5]:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	De
488	4472-LVYGI	Female	0	Yes	Yes	0	No	No phone service	DSL	Yes	...	
753	3115-CZMZD	Male	0	No	Yes	0	Yes	No	No	No internet service	...	
936	5709-LVOEQ	Female	0	Yes	Yes	0	Yes	No	DSL	Yes	...	
1082	4367-NUYAO	Male	0	Yes	Yes	0	Yes	Yes	No	No internet service	...	
1340	1371-DWPAZ	Female	0	Yes	Yes	0	No	No phone service	DSL	Yes	...	
3331	7644-OMVMY	Male	0	Yes	Yes	0	Yes	No	No	No internet service	...	
3826	3213-VVOLG	Male	0	Yes	Yes	0	Yes	Yes	No	No internet service	...	
4380	2520-SGTTA	Female	0	Yes	Yes	0	Yes	No	No	No internet service	...	
5218	2923-ARZLG	Male	0	Yes	Yes	0	Yes	No	No	No internet service	...	
6670	4075-WKNIU	Female	0	Yes	Yes	0	Yes	Yes	DSL	No	...	
6754	2775-SEFEE	Male	0	No	Yes	0	Yes	Yes	DSL	Yes	...	

11 rows × 21 columns

```
In [6]: df["TotalCharges"].replace(' ',np.nan, inplace=True)
df["TotalCharges"]=df["TotalCharges"].astype(float)
#Replacing the blank space with null value and converting the datatype to float
```


```
In [7]: df.describe()
```

Out[7]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7032.000000
mean	0.162147	32.371149	64.761692	2283.300441
std	0.368612	24.559481	30.090047	2266.771362
min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	401.450000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3794.737500
max	1.000000	72.000000	118.750000	8684.800000

```
In [8]: sns.distplot(df.TotalCharges)
```

Out[8]: <AxesSubplot:xlabel='TotalCharges', ylabel='Density'>



When the data is skewed, it is good to consider using the median value for replacing the null values. Since the above Distribution Plot signifies

the data is right skewed , replacing the null values with the median.

```
In [9]: df["TotalCharges"]=df["TotalCharges"].fillna(df["TotalCharges"].median())
```

```
In [10]: df.describe()
```

Out[10]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges
count	7043.000000	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692	2281.916928
std	0.368612	24.559481	30.090047	2265.270398
min	0.000000	0.000000	18.250000	18.800000
25%	0.000000	9.000000	35.500000	402.225000
50%	0.000000	29.000000	70.350000	1397.475000
75%	0.000000	55.000000	89.850000	3786.600000
max	1.000000	72.000000	118.750000	8684.800000

The datasets descriptions shows that the average Monthly Charges are USD 64.76 and it has maximum gone upto USD 118.75 The average tenure the telecom company has is of 32 months. The average Total Amount charged to the customer is USD 2283 and the maximum amount is USD 8684. By observation, I found that the MonthlyCharges and TotalCharges are not evenly distributed

Dropping the columns with least correlations

```
In [11]: df.drop(['customerID'],axis=1,inplace=True)
```

Dropping the costumerID column as it has the least correlation with the Churn column

Handling Categorical Data


```
In [12]: df_num=df.select_dtypes(['float64','int64'])
df_cat=df.select_dtypes(['object'])
```

Seperating the numerical and categorical columns

```
In [13]: df_cat.head()
```

Out[13]:

	gender	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	S
0	Female	Yes	No	No	No phone service	DSL	No	Yes	No	No	
1	Male	No	No	Yes	No	DSL	Yes	No	Yes	No	
2	Male	No	No	Yes	No	DSL	Yes	Yes	No	No	
3	Male	No	No	No	No phone service	DSL	Yes	No	Yes	Yes	
4	Female	No	No	Yes	No	Fiber optic	No	No	No	No	



```
In [14]: from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
```

Performing Label Encoding on the categorical columns so as to convert then into the machine readable form!

```
In [15]: for i in df_cat:
le=LabelEncoder()
df_cat[i]=le.fit_transform(df_cat[i])
```

```
In [16]: df_cat.head()
```

```
Out[16]:
```

	gender	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	S
0	0	1	0	0	1	0	0	2	0	0	
1	1	0	0	1	0	0	2	0	2	0	
2	1	0	0	1	0	0	2	2	0	0	
3	1	0	0	0	1	0	2	0	2	2	
4	0	0	0	1	0	1	0	0	0	0	



```
In [17]: df2=pd.concat([df_num,df_cat],axis=1)
```

```
In [18]: df2.head()
```



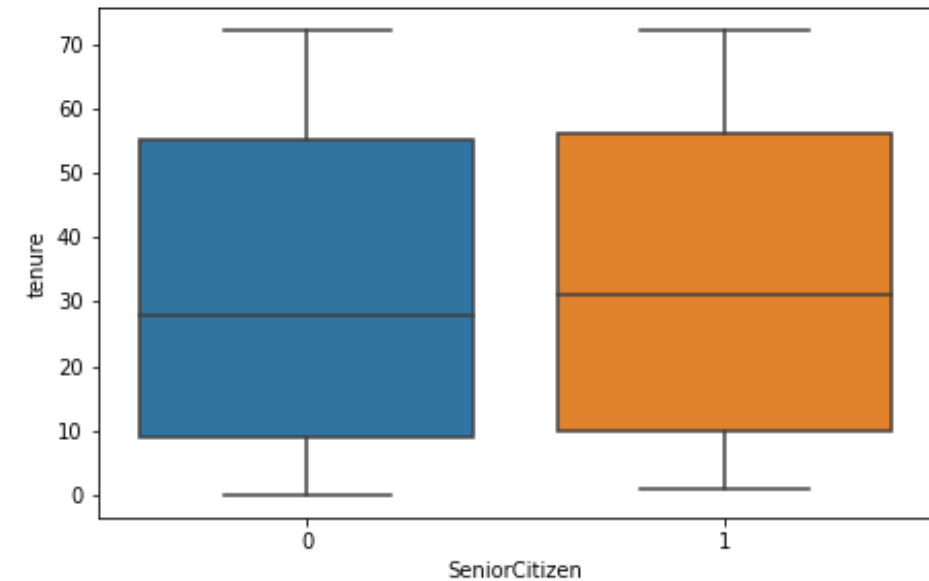

```
Out[18]:
```

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSecu
0	0	1	29.85	29.85	0	1	0	0	1	0	
1	0	34	56.95	1889.50	1	0	0	1	0	0	
2	0	2	53.85	108.15	1	0	0	1	0	0	
3	0	45	42.30	1840.75	1	0	0	0	1	0	
4	0	2	70.70	151.65	0	0	0	1	0	1	



Checking For Outliers


```
In [19]: fig,axes=plt.subplots(2,2,figsize=(16,10))
sns.boxplot(data=df, x="Contract", y="TotalCharges",ax=axes[0,0])
sns.boxplot(data=df2, x="SeniorCitizen" , y="tenure",ax=axes[0,1])
sns.boxplot(data=df2, x="Partner" , y="tenure",ax=axes[1,0])
sns.boxplot(data=df2, x="Churn" , y="TotalCharges",ax=axes[1,1])
plt.show()
```



Removing Outliers

```
In [20]: outlier1= df.loc[(df['Contract']=='Month-to-month') & (df['TotalCharges']>8000)]
outlier2= df.loc[(df['Churn']=='Yes') & (df['TotalCharges']>8500)]
print(outlier1.index,outlier2.index)
```

Int64Index([3820], dtype='int64') Int64Index([4610], dtype='int64')

```
In [21]: df2.drop(index=3820,inplace=True)
```

```
In [22]: df2.drop(index=4610,inplace=True)
```

By observing the boxplots , came across some outliers and removed them

```
In [23]: df2.head()
```

Out[23]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineSecu
0	0	1	29.85	29.85	0	1	0	0	1	0	
1	0	34	56.95	1889.50	1	0	0	1	0	0	
2	0	2	53.85	108.15	1	0	0	1	0	0	
3	0	45	42.30	1840.75	1	0	0	0	1	0	
4	0	2	70.70	151.65	0	0	0	1	0	1	

Reducing Skewness

```
In [24]: for i in df_num:
        print(f"{i}={df2[i].skew()}")
```

```
SeniorCitizen=1.8345916532810498
tenure=0.23995276839384197
MonthlyCharges=-0.22068712827085618
TotalCharges=0.963347739293753
```

TotalCharges column is highly skewed

```
In [25]: df2["TotalCharges"]=np.sqrt(df2["TotalCharges"])
```


```
In [26]: df2["TotalCharges"].skew()
```

```
Out[26]: 0.3096188474579813
```

Reduced the skewness





Data Exploration







```
In [27]: data=[df.Churn.value_counts()[1],df.Churn.value_counts()[0]]  
label=['Yes','No']  
plt.pie(data,labels=label, shadow=True, autopct="%0.1f%%",explode=[0,0.1],startangle=90)  
plt.show()
```





26.5% Customers have left within the last month.

```
In [28]: fig,axes=plt.subplots(4,2,figsize=(16,25))
sns.countplot(data=df,x='PhoneService',hue='Churn',ax=axes[0,0])
sns.countplot(data=df,x='SeniorCitizen',hue='Churn',ax=axes[0,1])
sns.countplot(data=df,x='InternetService',hue='Churn',ax=axes[1,0])
sns.countplot(data=df,x='gender',hue='Churn',ax=axes[1,1])
sns.countplot(data=df,x='Contract',hue='Churn',ax=axes[2,0])
sns.countplot(data=df,x='OnlineSecurity',hue='Churn',ax=axes[2,1])
sns.countplot(data=df,x='StreamingTV',hue='Churn',ax=axes[3,0])
sns.countplot(data=df,x='TechSupport',hue='Churn',ax=axes[3,1])
plt.show()
```

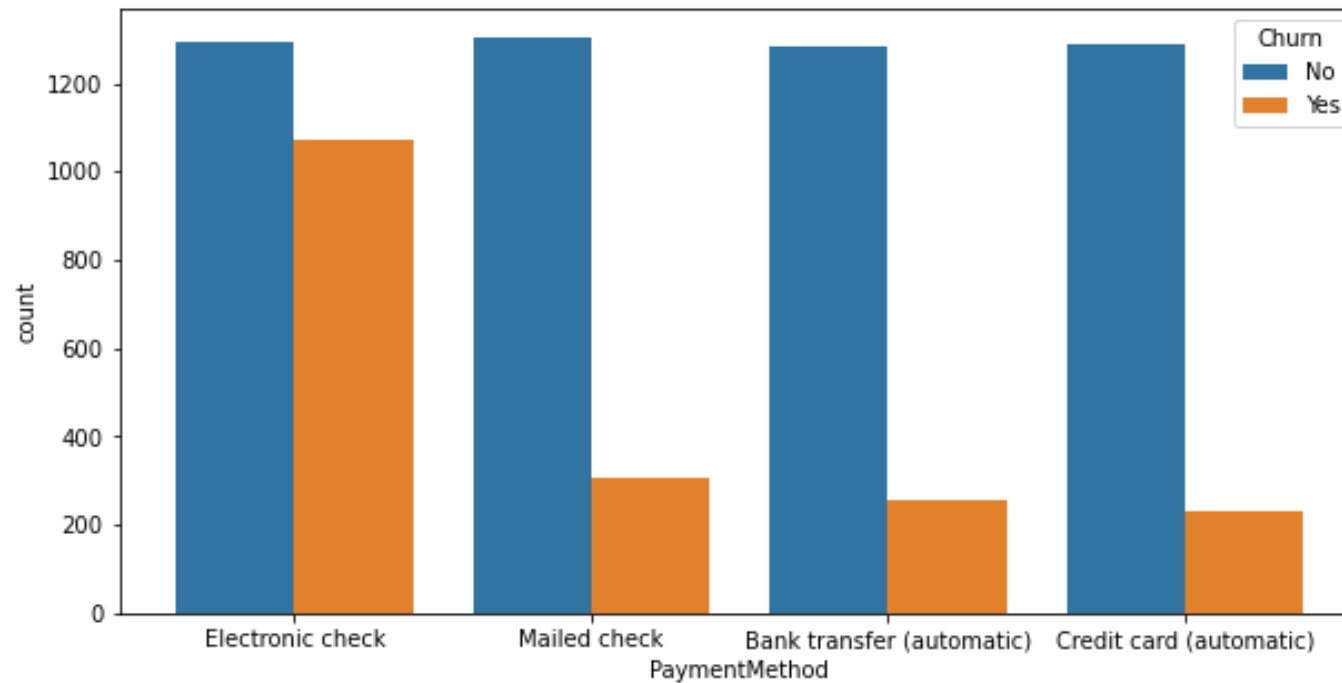






```
In [29]: fig, axes = plt.subplots(figsize=(10, 5))
sns.countplot(data=df, x='PaymentMethod', hue='Churn')
plt.plot()
```


Out[29]: []




The above countplots describe the Churn analysis with other features in the dataset

```
In [30]: sns.pairplot(data=df, vars=['MonthlyCharges', 'TotalCharges', 'tenure'], hue='Churn', palette='husl')
```

```
Out[30]: <seaborn.axisgrid.PairGrid at 0x213119b1d30>
```




```
In [31]: #Histogram of MonthlyCharges Frequency Distribution
plt.figure(figsize=(7,7))
plt.hist(df2["MonthlyCharges"], bins=10)
plt.xticks(rotation="vertical")
plt.title("MonthlyCharges Frequency Distribution")
plt.ylabel("No. of Customers")
plt.xlabel("MonthlyCharges")
plt.show()
```



Large number of customers (around 1600) pay MonthlyCharges between USD 20 to USD 30. The MonthlyCharges frequency is not evenly

distrubuted

```
In [32]: df2.corr()['Churn']
```

```
Out[32]: SeniorCitizen    0.151278
         tenure          -0.352681
         MonthlyCharges   0.193197
         TotalCharges     -0.225502
         gender           -0.008765
         Partner          -0.150660
         Dependents       -0.164169
         PhoneService      0.011895
         MultipleLines     0.037883
         InternetService  -0.047328
         OnlineSecurity    -0.289667
         OnlineBackup     -0.195799
         DeviceProtection -0.178401
         TechSupport       -0.282844
         StreamingTV       -0.036769
         StreamingMovies   -0.038679
         Contract         -0.396982
         PaperlessBilling  0.191761
         PaymentMethod     0.107407
         Churn             1.000000
         Name: Churn, dtype: float64
```

Correlation of the target variable with other features

According to the analysis, I came to the following conclusions:

1. 26.5% Customers have left the Company within last month.
2. The company has more no of Younger Generation than Senior Citizens with almost almost equal no of Male and Females.
3. Younger Generation Churn more than compared to the Senior Citizens. So the company must come up with more ideas in interest of the Young people.
4. The company provides phone services to many of the customers (i.e 6361 out of 7041).
5. The customers who is being provided Online Security Churn less. So the Company should increase the number of Customers with Online Security.

6. The customers with Technical Support tend to Churn less. So the Company should increase the number of Customers with Technical Support
7. The Customers with DSL as the Internet Service churn less as compared to the customers with Optical Fiber.
8. The Customers using Electronic Check as their Payment Method Churn more comparatively.
9. Customers with Yearly Contract tend to churn less as compared to customers with Monthly COntract

Scaling

```
In [33]: from sklearn.preprocessing import MinMaxScaler  
mm=MinMaxScaler()  
df2[["TotalCharges"]]=mm.fit_transform(df2[["TotalCharges"]])
```

Performed Feature Scaling on Total Charges

Seperating the Dependent and Independent variables for Training and Testing of Data

```
In [34]: x=df2.iloc[:, :-1]  
x
```

Out[34]:

	SeniorCitizen	tenure	MonthlyCharges	TotalCharges	gender	Partner	Dependents	PhoneService	MultipleLines	InternetService	OnlineS
0	0	1	29.85	0.012700	0	1	0	0	1	0	
1	0	34	56.95	0.440730	1	0	0	1	0	0	
2	0	2	53.85	0.068292	1	0	0	1	0	0	
3	0	45	42.30	0.434374	1	0	0	0	1	0	
4	0	2	70.70	0.089861	0	0	0	1	0	1	
...
7038	0	24	84.80	0.453644	1	1	1	1	2	0	
7039	0	72	103.20	0.917574	0	1	1	1	2	1	
7040	0	11	29.60	0.160798	0	1	1	0	1	0	
7041	1	4	74.40	0.148374	1	1	0	1	2	1	
7042	0	66	105.65	0.882932	1	0	0	1	0	1	

7041 rows × 19 columns



```
In [35]: y=df2.iloc[:, -1]  
y
```

```
Out[35]: 0      0  
1      0  
2      1  
3      0  
4      1  
      ..  
7038   0  
7039   0  
7040   0  
7041   1  
7042   0  
Name: Churn, Length: 7041, dtype: int32
```

```
In [36]: from sklearn.model_selection import train_test_split  
xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.25,random_state=1)
```

Model Building

Machine Learning Algorithms for Classification problems are: Logistic Regression, K Nearest Neighbor, Naive Bayes, Support Vector Machines (SVC)

```
In [37]: from sklearn.linear_model import LogisticRegression  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.naive_bayes import GaussianNB  
from sklearn.svm import SVC  
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

```
In [38]: def mymodel(model):  
        model.fit(xtest,ytest)  
        ypred=model.predict(xtest)  
        ac=accuracy_score(ytest,ypred)  
        cm=confusion_matrix(ytest,ypred)  
        cr=classification_report(ytest,ypred)  
        print(f" Accuracy: {ac} \n {cm} \n {cr} ")
```

```
In [39]: models=[]  
models.append(("Logistic ",LogisticRegression()))  
models.append(("KNN",KNeighborsClassifier()))  
models.append(("Naive bayes", GaussianNB()))  
models.append(("SVM", SVC()))
```

```
In [40]: for name,model in models:
          print(name)
          mymodel(model)
          print("\n\n")
```

Logistic

Accuracy: 0.8091993185689949

[[1198 122]

[214 227]]

	precision	recall	f1-score	support
0	0.85	0.91	0.88	1320
1	0.65	0.51	0.57	441
accuracy			0.81	1761
macro avg	0.75	0.71	0.73	1761
weighted avg	0.80	0.81	0.80	1761

KNN

Accuracy: 0.8387279954571266

[[1211 109]

[175 266]]

	precision	recall	f1-score	support
0	0.87	0.92	0.90	1320
1	0.71	0.60	0.65	441
accuracy			0.84	1761
macro avg	0.79	0.76	0.77	1761
weighted avg	0.83	0.84	0.83	1761

Naive bayes

Accuracy: 0.7654741624077229

[[1013 307]

[106 335]]

	precision	recall	f1-score	support
0	0.91	0.77	0.83	1320
1	0.52	0.76	0.62	441
accuracy			0.77	1761
macro avg	0.71	0.76	0.72	1761
weighted avg	0.81	0.77	0.78	1761

SVM

Accuracy: 0.7995457126632595

[[1233 87]

[266 175]]

	precision	recall	f1-score	support
0	0.82	0.93	0.87	1320
1	0.67	0.40	0.50	441
accuracy			0.80	1761
macro avg	0.75	0.67	0.69	1761
weighted avg	0.78	0.80	0.78	1761

Logistic Regression and KNN works well on the above dataset. However KNN gives more accuracy compared to other Machine Learning Algorithms with an accuracy of 83%